

```
In [1]: # Problem 3
import numpy as np
import matplotlib.pyplot as plt
import random

In [2]: # find the gradient of intercept
def intercept_gradient(y, y_predicted):
    gradient = (y_predicted - y)
    return gradient

In [3]: # find the gradient of weight
def weight_gradient(x, y, y_predicted):
    gradient= x*(y_predicted - y)
    return gradient

In [4]: # find the predicted value
def predict_y(x, weights, intercept):
    y_predicted = x @ weights + intercept
    return y_predicted

In [5]: # stochastic gradient descent
def SGD(X, Y):
    m = np. shape(X) [0]    # total number of samples
    n = np. shape(X) [1]    # total number of features
    weights = np. zeros(X. shape[1]) # initialize the weights vector
    intercept = 0           # initialize the intercept
    number_of_iterations = 5000  # total number of steps
    learning_rate = 0.01       # set the learning rate
    square_error = []
    random. seed(265)

    for i in range(number_of_iterations):
        r = random. randint(0, 14447)  # 14447 is the length of training set
        x = X[r]
        y = Y[r]
        y_predicted = predict_y(x, weights, intercept)
        error = y_predicted - y
        se = error**2
        square_error. append(se)

        # update weights
        weights = weights - learning_rate * weight_gradient(x, y, y_predicted)

        # update intercept
        intercept = intercept - learning_rate * intercept_gradient(y, y_predicted)

    # plot a graph of squared error versus number of updates
    plt. plot(np. arange(1, number_of_iterations), square_error[1:])
    plt. ylim((0, 1))
    plt. xlabel("number of updates")
    plt. ylabel("squared error")

    return weights, intercept

In [6]: # import the dataset
from sklearn. datasets import fetch_california_housing
california_housing = fetch_california_housing(as_frame=True)
DF = california_housing. frame
DF. head()
```

Out[6]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedHouseVal
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

```
In [7]: # b & c) choose the independent features as X and choose target feature Y
x = DF[['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population',
        'AveOccup', 'Latitude', 'Longitude']]. values
y = DF[['MedHouseVal']]. values
```


```
In [8]: # d) 0 - 1 normalization on X and Y
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x = scaler.fit_transform(x)
y = scaler.fit_transform(y)
```

```
In [9]: # randomly split data into training and test sets
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.3,random_state = 265)
```

```
In [10]: # g) get the best set of parameters
from sklearn.metrics import mean_squared_error
weights,intercept = SGD(x_train,y_train)
print("Weights:",weights)
print("Intercept:", intercept)

# use test set to calculate the MSE
y_predicted = []
for x in x_test:
    yhat = predict_y(x,weights,intercept)
    y_predicted.append(yhat)
print("MSE:",mean_squared_error(y_test,y_predicted))
```

Weights: [ 0.77752415 0.11892458 0.03479236 0.00530762 0.00893102 -0.00379474  
 -0.13080282 -0.07408918]  
 Intercept: [0.21669483]  
 MSE: 0.031076966697939037



```
In [11]: # h)
# The weights are [ 0.77752415  0.11892458  0.03479236  0.00530762
#                0.00893102 -0.00379474 -0.13080282 -0.07408918]
# The intercept are 0.21669483
# From the weights, we can find that MedInc which explain the house prices the most.
```

```
In [12]: # Problem 4
```

```
In [13]: # c) use sklearn SGDRegressor to find the best parameters and intercept
from sklearn.linear_model import SGDRegressor
sgd = SGDRegressor(max_iter=5000, alpha = 0.01, random_state = 265)
sgd.fit(x_train, y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
 y = column\_or\_1d(y, warn=True)

```
Out[13]: SGDRegressor(alpha=0.01, max_iter=5000, random_state=265)
```

```
In [14]: # d)
print(sgd.intercept_)
```

[0.27986693]

```
In [15]: print(sgd.coef_)

[ 0.74289115  0.12123235  0.02459013  0.00144189  0.00494056 -0.00390117
 -0.17298551 -0.13952611]
```

```
In [16]: # e)
# The weights are [ 0.74289115  0.12123235  0.02459013  0.00144189
#                  0.00494056 -0.00390117  -0.17298551 -0.13952611]
# The intercept are 0.27986693
# From the weights, we can find that MedInc which explain the house prices the most.
# There is only a small difference between those two results. The reason of difference might be that SGD Gressor has a criterion
# of convergence of cost function so it might stop training early, which might cause the small difference between two results.
```

```
In [17]: # Problem 5
```

```
In [18]: # compute the mean matrix of a matrix
def Mean(matrix):
    sum = 0
    row = matrix.shape[0]
    column = matrix.shape[1]
    for i in range(column):
        sum = 0
        for j in range(row):
            sum += matrix[j][i]
        mean = sum / row
        mean_column = np.full(row, mean)
        matrix[:,i] = mean_column
    return matrix
```

```
In [19]: # a function that used to perform dot product between two matrices
def Dot(matrix1, matrix2):
    output_matrix_list = []
    for i in range(matrix1.shape[0]):
        row_list = []
        for j in range(matrix2.shape[1]):
            number = sum([n*m for (n,m) in zip(matrix1[i,:],matrix2[:,j])])
            row_list.append(number)
        output_matrix_list.append(row_list)
    output_matrix = np.array(output_matrix_list)
    return output_matrix
```

```
In [20]: # a function that used to perform transpose of a matrix
def Transpose(matrix):
    row = matrix.shape[0]
    column = matrix.shape[1]
    transpose_matrix = np.zeros((column,row))
    for i in range(row):
        for j in range(column):
            transpose_matrix[j][i] = matrix[i][j]
    return transpose_matrix
```

```
In [21]: # test the function
A = DF[['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']].values
B = DF[['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']].values

mean = Mean(A)
cov = (Dot(Transpose(B - mean), (B - mean)))/(A.shape[0] - 1)
print(cov)

[[ 3.60932256e+00 -2.84614028e+00  1.53656801e+00 -5.58575949e-02
  1.04009792e+01  3.70288896e-01 -3.23859753e-01 -5.77647021e-02]
 [-2.84614028e+00  1.58396260e+02 -4.77288245e+00 -4.63718412e-01
 -4.22227058e+03  1.72429796e+00  3.00345508e-01 -2.72824366e+00]
 [ 1.53656801e+00 -4.77288245e+00  6.12153272e+00  9.93867801e-01
 -2.02333712e+02 -1.24688866e-01  5.62235473e-01 -1.36518371e-01]
 [-5.58575949e-02 -4.63718412e-01  9.93867801e-01  2.24591500e-01
 -3.55272253e+01 -3.04242537e-02  7.05752856e-02  1.26704371e-02]
 [ 1.04009792e+01 -4.22227058e+03 -2.02333712e+02 -3.55272253e+01
  1.28247046e+06  8.21712002e+02 -2.63137814e+02  2.26377839e+02]
 [ 3.70288896e-01  1.72429796e+00 -1.24688866e-01 -3.04242537e-02
  8.21712002e+02  1.07870026e+02  5.24916416e-02  5.15187178e-02]
 [-3.23859753e-01  3.00345508e-01  5.62235473e-01  7.05752856e-02
 -2.63137814e+02  5.24916416e-02  4.56229264e+00 -3.95705372e+00]
 [-5.77647021e-02 -2.72824366e+00 -1.36518371e-01  1.26704371e-02
  2.26377839e+02  5.15187178e-02 -3.95705372e+00  4.01413937e+00]]
```

```
In [22]: # Problem 7
```

```
In [23]: # a) perform sigmoid function
def sigmoid_f(z):
    return 1.0/(1.0 + np.exp(-z))
```

```
In [24]: # b) from result of sigmoid function to the predicted label
def classifier_f(x, w, b):
    m = x.shape[1]
    y_prediction = np.zeros((1,m))
    yh = sigmoid_f(np.dot(w.T, x)+b)    # get the result of sigmoid function

    # determine its label(0 or 1)
    for i in range(m):
        if yh[0, i] < 0.5:
            y_prediction[0,i] = 0
        else:
            y_prediction[0,i] = 1
    return y_prediction
```

```
In [25]: # c) cost function
def binary_loss_f(yh, y):
    m = y.shape[1]
    cost = -1/m * (np.sum(y * np.log(yh) + (1 - y) * np.log(1 - yh)))
    return cost
```

```
In [26]: # d) find the gradient of weights and intercept
def gradient_f(x, yh, y):
    m = x.shape[1]
    error = yh - y

    weight_gradient = 1/m * (np.dot(x, error.T))
    intercept_gradient = 1/m * (np.sum(error))

    # use dictionary to store gradients
    grads = {"weight": weight_gradient, "intercept": intercept_gradient}

    return grads
```

```
In [27]: # e)
def optimizer_f(x, y):
    n = x.shape[0]
    weights = np.zeros((n,1))
    intercept = 0
    learning_rate = 0.05
    epochs = 10000
    cost = []

    for i in range(epochs):
        yh = sigmoid_f(np.dot(weights.T, x)+intercept)
        grads = gradient_f(x, yh, y)

        # update weights
        weights = weights - learning_rate*grads["weight"]
        # update intercept
        intercept = intercept - learning_rate * grads["intercept"]

        cost.append(binary_loss_f(yh, y))

    # find the predicted value of y to calculate the accuracy
    y_predicted = classifier_f(x, weights, intercept)
    accuracy = 1 - np.mean(np.abs(y_predicted - y))
    print("Accuracy", accuracy)

    # plot the graph of cost versus number of updates
    plt.plot(np.arange(1, epochs), cost[1:])
    plt.xlabel("number of updates")
    plt.ylabel("cost")

    return weights, intercept
```

```
In [28]: # test the functions
from sklearn.datasets import load_breast_cancer
breast_cancer = load_breast_cancer(as_frame = True)
DF = breast_cancer.frame
DF.head()
```

Out[28]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	...	worst texture	worst perimeter	worst area	worst smoothness
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419	0.07871	...	17.33	184.60	2019.0	0.1
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812	0.05667	...	23.41	158.80	1956.0	0.1
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069	0.05999	...	25.53	152.50	1709.0	0.1
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597	0.09744	...	26.50	98.87	567.7	0.2
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809	0.05883	...	16.67	152.20	1575.0	0.1

5 rows × 31 columns




```
In [29]: # a & b)
x = DF.iloc[:, :30].values
y = DF[['target']].values
```

```
In [30]: # c)
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x = scaler.fit_transform(x)
y = scaler.fit_transform(y)
x = x.T
y = y.T
```

```
In [31]: # d)
weights, intercept = optimizer_f(x, y)
print("Weights:", weights)
print("Intercept:", intercept)
```

Accuracy 0.968365553602812  
Weights:  $\begin{bmatrix} -1.1434992 & -1.30912899 & -1.24649986 & -1.57207726 & 0.4750165 & -0.71928016 & -2.48155437 & -3.34113606 & 0.41115119 & 2.03552652 & -1.64607791 & 0.46426397 & -1.29014345 & -1.22181241 & 0.5989249 & 0.91627859 & 0.51280409 & 0.27778245 & 0.7602846 & 0.93372269 & -2.3582086 & -2.14580201 & -2.21834932 & -2.15158438 & -0.98440036 & -1.12134051 & -1.78357458 & -3.27399809 & -1.05564681 & -0.1960431 \end{bmatrix}$   
Intercept: 7.596854589203279



```
In [32]: # e) report Final equatoin
names = DF.columns.values.tolist()
print("The final function is:",intercept, end = "")
for i in range(len(weights)-1):
    print(" %f * %s +"%(weights[i],names[i]),end = "")
print(" %f * %s"%(weights[len(weights)-1],names[len(weights)-1]),end = "")
```

The final function is: 7.596854589203279 -1.143499 \* mean radius + -1.309129 \* mean texture + -1.246500 \* mean perimeter + -1.572077 \* mean area + 0.475017 \* mean smoothness + -0.719280 \* mean compactness + -2.481554 \* mean concavity + -3.341136 \* mean concave points + 0.411151 \* mean symmetry + 2.035527 \* mean fractal dimension + -1.646078 \* radius error + 0.464264 \* texture error + -1.290143 \* perimeter error + -1.221812 \* area error + 0.598925 \* smoothness error + 0.916279 \* compactness error + 0.512804 \* concavity error + 0.277782 \* concave points error + 0.760285 \* symmetry error + 0.933723 \* fractal dimension error + -2.358209 \* worst radius + -2.145802 \* worst texture + -2.218349 \* worst perimeter + -2.151584 \* worst area + -0.984400 \* worst smoothness + -1.121341 \* worst compactness + -1.783575 \* worst concavity + -3.273998 \* worst concave points + -1.055647 \* worst symmetry + -0.196043 \* worst fractal dimension

```
In [33]: # f)
# Rank parameters from positive to negative (biggest to smallest)
# 2.03552652 mean fractal dimension
# 0.93372269 fractal dimension error
# 0.91627859 compactness error
# 0.7602846 symmetry error
# 0.5989249 smoothness error
# 0.51280409 concavity error
# 0.4750165 mean smoothness
# 0.46426397 texture error
# 0.41115119 mean symmetry
# 0.27778245 concave points error
# -0.1960431 worst fractal dimension
# -0.71928016 mean compactness
# -0.98440036 worst smoothness
# -1.05564681 worst symmetry
# -1.12134051 worst compactness
# -1.1434992 mean radius
# -1.22181241 area error
# -1.24649986 mean perimeter
# -1.29014345 perimeter error
# -1.30912899 mean texture
# -1.57207726 mean area
# -1.64607791 radius error
# -1.78357458 worst concavity
# -2.14580201 worst texture
# -2.15158438 worst area
# -2.21834932 worst perimeter
# -2.3582086 worst radius
# -2.48155437 mean concavity
# -3.27399809 worst concave points
# -3.34113606 mean concave points

# We can see that there are more coefficients that are more negatively related coefficient to the target than
# positively related ones.
# Since we have normalized the data, we should see the absolute value of the coefficients
# the bigger the absolute value, the more impact on the target.
# Top 3 impact on the target are C8, C28, and C7, all negatively related.
# And the top 3 positively related are C10, C20, C16
```

```
In [34]: # Problem 10
```

```
In [35]: # import dataset
from sklearn.datasets import fetch_california_housing
california_housing = fetch_california_housing(as_frame=True)
DF = california_housing.frame
DF.head()
```

Out[35]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedHouseVal
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23	4.526
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22	3.585
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24	3.521
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25	3.413
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25	3.422

```
In [36]: # a & b)
x = DF[['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population',
        'AveOccup', 'Latitude', 'Longitude']].values
y = DF[['MedHouseVal']].values
```



```
In [37]: # c)
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x = scaler.fit_transform(x)
y = scaler.fit_transform(y)
```

```
In [44]: # apply LeaveOneOut cross-validation and calculate MSE
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import LeaveOneOut
from sklearn.metrics import mean_squared_error
cv = LeaveOneOut()
model = LinearRegression()
y_true = []
y_pred = []
for i, j in cv.split(x):
    x_train = x[i, :]
    x_test = x[j, :]
    y_train = y[i]
    y_test = y[j]
    model.fit(x_train, y_train)
    y_predicted = model.predict(x_test)
    y_true.append(y_test[0][0])
    y_pred.append(y_predicted[0][0])
print(mean_squared_error(y_true, y_pred))
```

20640

0.02245687523556031

```
In [45]: # apply KFold cross-validation and calculate MSE
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import KFold
from sklearn.metrics import mean_squared_error
cv = KFold(shuffle = True, random_state = 265)
model = LinearRegression()
y_true = []
y_pred = []
for i, j in cv.split(x):
    x_train = x[i, :]
    x_test = x[j, :]
    y_train = y[i]
    y_test = y[j]
    model.fit(x_train, y_train)
    y_predicted = model.predict(x_test)
    y_true.append(y_test[0][0])
    y_pred.append(y_predicted[0][0])
print(mean_squared_error(y_true, y_pred))
```

5

0.01065684376084764

```
In [40]: # apply Train-test split cross-validation and calculate MSE
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
model = LinearRegression()
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3, random_state = 265)
model.fit(x_train, y_train)
y_predicted = model.predict(x_test)
print(mean_squared_error(y_test, y_predicted))
```

0.023161877546173583

```
In [41]: # The MSE of L00CV is 0.0224569
# The MSE of KFold is 0.0106568
# The MSE of Train-test split is 0.0231619
# We can find that the MSE of KFold is the smallest, and MSE of other two is nearly same.
# The reason of difference between L00CV and KFold might be that the number of samples to calculate MSE of KFold is only k = 5,
# but the number of samples to calculate MSE of L00CV is 20640, which will increase MSE.
# And the reason of difference between L00CV, KFOLD and Train-test split is that there are more training data for L00CV and KFold.
```